

Pediment Bank · 3-Minute Demo Runbook

Snowflake Summit · Pediment Booth · for the recorded video walkthrough of the Pediment Bank dbt-wizard demo. Looped unattended on the booth screen.

3-5 days

MANUAL MODEL BUILD TODAY

90 sec

DBT-WIZARD BUILD

\$2.4M

CNP FRAUD BLOCKED / WK

8 hr

EARLIER THAN RISK + FRAUD DAILY STANDUP

ODI architecture · five layers, one loop

Fivetran

Connectors + Custom SDK

MDLS

Managed Data Lake (Iceberg on S3)

dbt Labs

Bronze · Silver · Gold marts

dbt-wizard

Build-time AI · 90s authoring

Snowflake

Reads governed gold layer

DATA ENGINEERS ANCHOR HERE

Architecture · Wizard · Live · Pipeline

Show the loop, the four sub-agents, the materialization, the connector deep-links. This is your DE proof.

BUSINESS USERS ANCHOR HERE

Deposits, Lending, Fraud, AML, Commercial

Show the answers, the KPI tiles, the segments, the executive read. The same governed gold layer powers all of it.

Frame the demo on business outcomes — never functionality

PAIN

The fraud team can't size new spikes before they bleed

REASON

Every new question = 3-5 day DE backlog

IMPACT

Realized fraud losses, OCC findings, blocked-customer false positives

CAPABILITY

Build-time AI on the governed gold layer

PROOF

Live model build in 90 seconds — on screen

The mantra: the win is not that dbt-wizard *can* author a model. The win is that the Fraud Operations Director walks into the next Risk + Fraud daily standup with a tested, lineage-tracked answer that did not exist this morning. Every beat below ties to that one sentence.

The 3-minute script · five beats

00:00 – 00:25 Beat 1 · The stakeholder hook

WHAT YOU SAY (25S)

"It is Friday morning at Pediment Bank. The Fraud Operations Director needs to know: Why did CNP fraud spike for tier-2 merchants in Texas last Tuesday? The Risk + Fraud daily standup happens in 8 hours. Today, an answer like this costs three-to-five days of data-engineering work and a missed meeting. We are about to change that."

WHAT YOU SHOW

Open the **Scenario page** (/scenario). The Fraud Operations Director's question is the first thing in frame. Let the T-minus counter tick.

OUTCOME → Anchors the audience on a real business question, not a feature.

00:25 – 00:50 Beat 2 · The cost of the gap

WHAT YOU SAY (25S)

"There is no gold-layer mart that answers this. Three days from now, after a request, a backlog, a scope, a build, and a peer review, the engineering team might land it. The Risk + Fraud daily standup will have already happened. A 2.4x tier-2 spike in Texas — every quiet day means about \$2.4 million in card-not-present exposure that nobody is acting on — and that is just the dollars we can see."

WHAT YOU SHOW

Stay on Scenario. Highlight the four KPI tiles: 3–5 days, 90s, \$2.4M, T-minus to Risk + Fraud daily standup.

OUTCOME → Quantifies the cost of the gap. Sets up the 90-second proof.

00:50 – 02:20 Beat 3 · The live build (proof moment · DE-anchor)

WHAT YOU SAY (90S)

"We click Run the Build. Four sub-agents — Explorer, Summary, Worker, Verification — open the dbt project, find the right grain, pull a slice of the gold layer, write the new model, write its tests, and materialize it. The viewer sees the SQL appear character by character. By the time we are done talking, the model exists, the tests pass, the lineage graph has updated."

While narrating, pause once at: "Worker is anchoring the CTE to the grain the Fraud Operations Director asked for — not a notebook-grade approximation." Then again at: "Verification is writing a uniqueness test on the natural key. This is a governed gold asset, not a SQL snippet."

WHAT YOU SHOW

Click into [/wizard-live](#). Let the playback run at 1x. The agents type SQL and YAML in real time. The step rail moves Discovery → Schema → Inspection → Model Creation → Tests → Materialization. Terminal-window aesthetic.

OUTCOME → Data engineers see the loop end-to-end. The wow is that the output is a tested, governed asset — not a snippet pasted into Slack.

02:20 – 02:45 Beat 4 · The outcome (P&L tie-back · business-anchor)

WHAT YOU SAY (25S)

"The new mart is materialized. The Risk + Fraud daily standup convenes in 8 hours, not three days. The Fraud Operations Director walks in with the answer: three Houston-area merchant terminals chained to a single device-fingerprint cluster originated 67% of the spike — block at the terminal, not the cardholder. That is a merchant-side block, not a customer-friction change. And the next time anyone asks this question — card-present, ACH, wire, anywhere — the model is already there."

WHAT YOU SHOW

Cut to [/outcome](#). Materialized model, tests passed, lineage updated, root-cause panel, impact panel. Then quick scroll through the business-user pages (Deposits, Lending, Fraud, AML, Commercial) so the audience sees the same gold layer powering the executive read.

OUTCOME → Closes the loop from question to action AND shows the business-user surface running on the same data.

02:45 – 03:00 Beat 5 · The architectural ask

WHAT YOU SAY (15S)

"Fivetran lands every source. MDLS owns the open lake. dbt governs the layer. dbt-wizard authors the gaps. Snowflake reads it. That loop closes for fraud rings, AML alert triage, deposit-beta sensitivity, and OCC audit prep — all on the same governed lake. Want to see how this maps to your gold layer? Stop by the Pediment booth or scan the QR."

WHAT YOU SHOW

Final card with the five-pillar loop, QR code to [fivetran.com/contact](#). Tag dbt Labs co-brand.

OUTCOME → Calls out the architectural pattern. Invites a discovery conversation.

Business outcomes to weave through — never lead with features

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT
Risk + Fraud daily standup gets the answer same-day	3–5 day DE backlog	"In time for the meeting, not three days late"

OUTCOME THE BUYER FEELS	WHAT IT REPLACES	HOW TO PHRASE IT
\$2.4M exposure sized in 90 seconds	Open gap window	"The cost of the gap closed before lunch"
Reusable governed mart for every future ask	One-off notebook scripts	"The next person who asks gets it for free"
CDO does not add headcount to keep up	Hiring DE backfill	"Your team stops being the bottleneck"
Auditor sees lineage + tests, not a query history	Manual evidence packs	"Examiner-ready by default"

Anti-patterns · what NOT to say

- **Don't say:** "It runs dbt show against the warehouse and authors a new model file." **Say:** "Ninety seconds from question to materialized answer."
- **Don't say:** "Cortex Analyst." Cortex is intentionally absent from this demo. The story is build-time AI on the governed lake.
- **Don't say:** "dbt-wizard generates SQL." Say: "dbt-wizard authors a governed asset — schema contract, tests, lineage, ownership."
- **Don't say:** "Look at the typing animation." The animation is the texture, not the value.
- **Don't open with the architecture diagram.** Open with the Fraud Operations Director's question. Architecture is the closer.
- **Don't skip the business pages.** After the wizard finishes, scroll through Deposits, Lending, Fraud, AML, Commercial so the audience sees the persona surface that runs on the same gold layer.

Discovery prompts to anchor before the demo (if presenting live)

BEFORE YOU PRESS PLAY

- "What questions does your Fraud Operations Director ask that your team cannot answer same-day?"
- "When a stakeholder asks for a new cut of data, what is the average time from ask to merged dbt model?"
- "How many one-off SQL scripts are sitting in someone's Snowflake history that never made it back into the governed layer?"
- "Who is your Economic Buyer for new data investment? Likely audience: CRO, Head of Fraud, CISO, CDO."

AFTER THE DEMO LANDS

- "Which of those agent steps — Discovery, Schema, Inspection, Model Creation, Tests, Materialization — is where your team spends the most time today?"
- "If we sized a pilot on three of your highest-value recurring Fraud Operations Director questions, what would success look like in 90 days?"
- "Who needs to be in a 30-minute architecture review to see how this lands on your gold layer? Audience: CRO, Head of Fraud, CISO, CDO."

Personas at the booth — split the time

DATA ENGINEER · 60% OF BOOTH TIME

Scenario → Live → Outcome → Architecture → Pipeline

Lead with the question, show the four sub-agents, then walk the architecture (Fivetran + MDLS + dbt Labs + dbt-wizard + Snowflake) and the Pipeline page with real Fivetran connector deep-links.

BUSINESS USER · 40% OF BOOTH TIME

Deposits, Lending, Fraud, AML, Commercial

After the wizard finishes, drop into the dashboard pages. Show that the same governed gold layer powers the CRO, Treasurer, CFO, Head of Commercial Banking read. The build-time AI is invisible to them — they see the answer.

Required follow-up materials

- One-pager · Pediment Bank · at fivetran-jasonchletsos.github.io/Fivetran-Demo-Repository/one-pager/banking/
- Live demo URL · fivetran-jasonchletsos.github.io/Banking-ODI-Demo/
- 3-minute walkthrough video (this recording)
- Architecture diagram on the Architecture tab of the demo
- Cross-industry credibility · Clarity Health

Booth-floor reminder: the demo is built to run unattended on the booth screen. If a prospect walks up mid-playback, do not restart. Pick up the narration where the playback is — that is also the point. The build finishes whether someone is watching or not. That is the entire pitch.